EECS 368/468 – Lab 1
Programming Massively Parallel Processors with CUDA

Matrix Multiplication in CUDA

In this first assignment, you will work in teams of two to write a matrix multiplication code in CUDA. Submissions from teams with only 1 or ≥3 members will automatically get zero points, unless the instructor explicitly allowed such team through written communication. If you need to find a partner, go to Piazza (our Canvas site has a link) and post a request.

The lab is due on Monday Jan 25 at 11:59pm. The late submission policies are explained in the syllabus. Please follow the instructions below to do the lab.

Install the CUDA SDK at your Wilkinson Lab account

1. Login to a machine at Wilkinson Lab (Tech M338). If you wish to work remotely (e.g., via ssh), the Wilkinson Lab hostnames are:

alfred.eecs.northwestern.edu	joker.eecs.northwestern.edu
bane.eecs.northwestern.edu	killercroc.eecs.northwestern.edu
batgirl.eecs.northwestern.edu	madhatter.eecs.northwestern.edu
batman.eecs.northwestern.edu	nightwing.eecs.northwestern.edu
clayface.eecs.northwestern.edu	poisonivy.eecs.northwestern.edu
cobblepott.eecs.northwestern.edu	ras.eecs.northwestern.edu
freeze.eecs.northwestern.edu	riddler.eecs.northwestern.edu
gordon.eecs.northwestern.edu	robin.eecs.northwestern.edu
gotham.eecs.northwestern.edu	scarecrow.eecs.northwestern.edu
harley.eecs.northwestern.edu	selina.eecs.northwestern.edu
huntress.eecs.northwestern.edu	twoface.eecs.northwestern.edu
hush.eecs.northwestern.edu	

2. Setup the CUDA environment. If you use csh or tcsh:
```
source /usr/local/cuda-5.0/cuda-env.csh
```
or if you use the bash shell:
```
. /usr/local/cuda-5.0/cuda-env.sh
```

3. Install SDK in your home directory
```
cp -r /home/hardav/NVIDIA_CUDA-5.0_Samples ~/;
cd ~/NVIDIA_CUDA-5.0_Samples;
make
```

4. Test by finding out the characteristics of your GPU device
```
cd ~/NVIDIA_CUDA-5.0_Samples/1_Utilities/deviceQuery;
./deviceQuery
```

Install the labs scaffold at your Wilkinson Lab account

5. Download the labs-scaffold tarball from Canvas into some directory in your Wilkinson Lab account.

6. Untar the code scaffold tarball:
```
tar xvf EECS468-CUDA-Labs-Scaffold.tgz
```

Install lab1 at your Wilkinson Lab account

7. Download the lab1 tarball from Canvas into some directory in your Wilkinson Lab account.

8. Install the lab1 tarball and compile the lab sources:
```
cd EECS468-CUDA-Labs
tar xvf EECS468-CUDA-Lab1.tgz
make
```

Complete the assignment

9. **Remember**: every time you login to the Wilkinson lab to program in CUDA you need to setup the CUDA environment. If you use csh or tcsh:
```
source /usr/local/cuda-5.0/cuda-env.csh
```
or if you use the bash shell:
```
. /usr/local/cuda-5.0/cuda-env.sh
```

10. **Assignment Task:**
Edit the following two functions
`MatrixMulOnDevice(...)` in the source file `matrixmul.cu`
`MatrixMulKernel(...)` in the source file `matrixmul_kernel.cu`
to implement matrix multiplication on the GPU device. **Do not change any other piece of the source code**. The size of the matrix is defined such that one thread block will be sufficient to compute the entire solution matrix.

The source code you will be working on is at :
```
EECS468-CUDA-Labs/labs/src/lab1
```

Compiling your code produces the executable:
```
EECS468-CUDA-Labs/labs/bin/linux/release/lab1
```

11. There are several modes of operation for the application.

    a. No arguments: the application will create two randomly initialized matrices to multiply. After the device multiplication is invoked, it will compute the correct solution matrix using the CPU, and compare that solution with the device-computed solution. If it matches (within a certain tolerance), it will print out "`Test PASSED`" to the screen before exiting.

    b. One argument: the application will use the random initialization to create the input matrices, and write the device-computed output to the file specified by the argument.

2

c. Two arguments: the application will initialize the two input matrices with the values found in the files provided as arguments. No output is written to file.

d. Three arguments: the application will read its inputs from the files provided by the first two arguments, and write its output to the file provided in the third.

Note that if you wish to use the output of one run of the application as an input, you must delete the first line in the output file, which displays the accuracy of the values within the file. The value is not relevant for this application.

12. Hand in your solution:

a. Create a tarball of your solution, which includes the contents of the `lab1` source folder provided, with all the changes and additions you made to the source files. Please make sure you do a `make clobber` before submitting; don't include object code or executables, as these are typically large files:

```
cd EECS468-CUDA-Labs/labs/src/lab1;
make clobber
tar cvfz EECS468-lab1-YourNames.tgz *
```

where you should replace `YourNames` with the **concatenated last names of the team members**.

b. Submit your solution tarball via Canvas. Only one student in a team needs to submit. If you submit multiple times, only the latest submission will be graded. Please do all submissions from the same account.

Good luck!